Peter A. McKee
Round Rock, TX 78681
peter.mckee@gmail.com

SUMMARY

Hands-on technical leader with deep, current experience architecting and shipping production AI systems—autonomous agent platforms, multi-tenant voice AI with RAG, and provider-agnostic LLM tooling. Three decades of engineering experience spanning IC, architect, and director-level roles at Docker, JFrog, AssemblyAI, and Sonar, leading global teams and turning ambiguous requirements into shipped product. Combines engineering depth with clear technical direction and a bias for building.

TECHNICAL SKILLS

Languages & Runtimes: Rust, TypeScript, JavaScript, Node.js, Python, Java, .NET
AI / ML: Anthropic Claude, OpenAI, Ollama, LLM integration, Prompt engineering, RAG, pgvector embeddings, Agentic workflows, Tool use / function calling, Multi-provider LLM abstraction, VAPI

Web & Frontend: React, Next.js, Redux, Angular, Tailwind v4

Backend & APIs: Express, Restify, REST / JSON APIs, NextAuth v5, OAuth (Google Calendar, Microsoft Graph), Webhooks, RESP wire protocol

Databases & Storage: PostgreSQL, Prisma, Redis, MongoDB, Elasticsearch, pgvector

Cloud & Serverless: AWS (Lambda, API Gateway, serverless), GCP, Docker, Kubernetes, Docker Swarm, Terraform

DevOps & Source Control: CI/CD, Microservices, GitHub, GitLab

Workflow Automation: n8n, Make, Zapier

Systems & Concurrency: Tokio async runtime, shared-nothing multi-shard architecture, extendible hashing

Tooling & Testing: Vitest, Puppeteer, Lightpanda, Tavily, pnpm / npm workspaces, Commander

Architecture: Multi-tenant SaaS, RBAC, Monorepos, Event-driven systems, Dependency injection

PROFESSIONAL EXPERIENCE

2/2026 - Present | ShrikeDB - Open Source Project (https://github.com/shrikedb/shrikedb)
Lead Developer & Architect (Rust, Tokio, Redis/RESP)

- Architected and built a Redis-compatible in-memory database in Rust, implementing ~95 commands across 6 data-type families (strings, keys, lists, hashes, sets, sorted sets) over the RESP wire protocol — drop-in compatible with redis-cli, redis-benchmark, and existing client libraries.
- Designed a shared-nothing, multi-shard execution model: each shard runs on its own OS thread with a dedicated single-threaded Tokio runtime and exclusively owns its slice of the keyspace, eliminating locks on the hot path and scaling throughput linearly with core count.
- Implemented a Dash-style extendible hash table (segments, fingerprint filtering, incremental directory doubling) that avoids stop-the-world rehashes and keeps tail latency flat under load; built a 16-byte CompactObj representation to inline short strings, integers, and small collections without heap allocation.
- Built CRC32 + hash-tag key routing with parallel fan-out and gather for multi-key commands (MGET, MSET, DEL, EXISTS), and shipped RDB persistence (save/load) coordinated across shards via an mpsc-routed coordinator task.

1/2025-Present | Ronin AI
Lead Developer, Technical Delivery

Lead a boutique AI consultancy, wearing multiple hats across developer experience, technical delivery, marketing, and hands-on implementation—operating with startup agility and minimal structure.

Selected Projects

Voice Assistant — Multi-tenant AI Phone Receptionist Platform

Next.js 16 · TypeScript · Prisma/Postgres · NextAuth v5 · Tailwind v4 · VAPI · Google Calendar / Microsoft Graph

SaaS platform that lets businesses stand up an AI phone receptionist on their own VAPI account in minutes — answers inbound calls, books appointments against the company's live calendar, and grounds responses in a per-tenant knowledge base, with full call history and role-based access.

- Designed and built a multi-tenant SaaS in Next.js 16 + TypeScript with row-level tenancy in Postgres/Prisma and a three-tier RBAC model (super-admin / company-admin / company-user) on NextAuth v5, scoping every query to the authenticated session's company so calls, appointments, knowledge, and credentials stay isolated across tenants.
- Integrated the VAPI telephony API to programmatically provision per-tenant AI assistants on a BYO-credentials model (each company connects their own VAPI account rather than the platform reselling minutes), with a secured webhook that ingests live call events and persists status, transcripts, summaries, and recordings into a per-tenant call log with playback.
- Shipped two production AI capabilities on top of the voice assistant: real-time appointment booking via VAPI tool-calls backed by Google Calendar and Microsoft Graph OAuth (availability checks, conflict detection, timezone handling, SMS confirmation), and a per-tenant knowledge base using chunked ingestion + pgvector embeddings so the assistant answers freeform business questions grounded in the company's own content.

Multi-Surface Agent Platform

Node.js · Next.js 16 · React 19 · Tailwind v4 · TypeScript · Anthropic / OpenAI / Ollama SDKs · Puppeteer + Lightpanda · Tavily · Redis · Vitest · npm workspaces

Built a full-stack autonomous AI agent platform as a Node.js monorepo (CLI + Next.js web dashboard + reusable core), supporting multiple LLM providers and headless browser automation for long-running, tool-using agent jobs.

- Designed a provider-agnostic core (`@agent/core`) with pluggable adapters for Anthropic, OpenAI, and Ollama, plus a tool registry, planner, context manager, working/long-term memory (Redis-backed), and per-tool permission gating.
- Shipped a Next.js 16 / React 19 web dashboard with streaming event logs, approval modals for human-in-the-loop tool calls, and per-job views, sharing the same agent runtime as the CLI.
- Integrated headless browser tooling (Lightpanda + Puppeteer) and Tavily web search to give agents reliable web-research and page-interaction capabilities under a unified tool interface.
- Architected the system as a pnpm/npm workspaces monorepo (`core`, `cli`, `web`) with dependency injection and event-driven observability, enabling Vitest unit coverage of the agent loop in isolation from UI.

AI Agent Framework (Spec-Driven)

Node.js · JavaScript + JSDoc typedefs · Anthropic SDK · Commander · @inquirer/prompts · fast-glob · js-yaml · chalk · ora · Vitest · pnpm

Designed and implemented a modular Node.js framework for autonomous planning agents that decompose a goal, generate a plan, execute via tools, evaluate results, and replan — built from a complete written spec set before any code.

- Authored a 6-module architecture (Core Orchestrator, Planning Engine, Memory, Tools, LLM Abstraction, Observability) with explicit dependency rules and per-module specs in `docs/`, enforcing a clean acyclic graph and dependency injection throughout.
- Implemented goal decomposition, plan generation, step evaluation, and replanning as separable strategies (iterative, full-plan, reactive), each driven by structured LLM output with runtime schema validation.

- Built a tool system with JSON-Schema input/output validation, configurable timeouts, retry policies, and human-approval gating, plus built-in `think`, `ask-user`, and `sub-task` primitives.
- Delivered usage examples (simple task, custom tools, human-in-the-loop, long-term memory) and a Vitest suite, with JSDoc `@typedef` contracts standing in for TypeScript across the codebase.

8/2024-12/2024 | AssemblyAI
Global Head of Community, Developer Relations, Developer Experience

- Led global DevRel + DX function spanning developer education, evangelism, documentation, and multi-language SDK strategy, aligning technical assets to onboarding milestones and developer success metrics.
- Owned the roadmap for SDK + documentation improvements across 6 languages, prioritizing friction reduction, consistent API patterns, and high-signal examples that accelerate "first successful call" and production readiness.
- Built a scalable technical asset pipeline—sample apps, demos, tutorials, videos, blogs, and guides—with operational processes to ship content globally and keep it aligned with product releases.
- Acted as the voice of the developer internally: gathered community feedback, synthesized integration pain points, and partnered with Product/Engineering to drive changes in developer workflows and tooling.
- Grew developer community and reach through high-impact education, including a YouTube channel exceeding 155K subscribers and 13M views, demonstrating the ability to scale developer trust and advocacy.
- Led and coached a team of six Developer Educators/Evangelists, defining messaging, technical tone, and content standards to make advanced topics approachable for developers at any level.

1/2023-8/2024 | Sonar
Vice President - Community & Developer Relations

- Served as executive owner for developer evangelism strategy across major language ecosystems; built programs that drive adoption through integration guidance, best practices, and developer-first storytelling.
- Led a global org (Developer Advocates + Community leadership) with a shared vision and roadmap for technical assets, field enablement, and community growth at scale.
- Created a framework for consistent, high-quality technical output (demos, reference apps, tutorials, keynotes, and blogs), ensuring content mapped to product capabilities and real developer workflows.
- Built tight cross-functional partnerships with Product/Engineering to translate developer feedback into product and docs improvements, prioritizing onboarding friction and integration reliability.
- Delivered thought leadership and technical talks to large audiences (reported reach 8M+ engineers/architects/technical leaders), strengthening credibility and accelerating community-led advocacy.
- Established measurable operating rhythms (program planning, launch readiness, KPIs) to scale global developer impact without sacrificing technical depth or voice.

5/2022-1/2023 | JFrog
Sr. Director - Developer Advocacy
As head of Developer Advocacy at JFrog, I led a global team to drive developer engagement through education, strategic partnerships, and community outreach. I represented JFrog at major events, creating demos and tutorials to help developers onboard and succeed with the platform.

5/2017-5/2022 | Docker, Inc.
Director - Head of Developer Relations & Community
3/2020 - 5/2022

- Owned global Developer Evangelism + Developer Experience strategy for Docker's developer ecosystem (15M+ developers), aligning community programs, technical assets, and messaging to improve onboarding and day-to-day product success.
- Served as the voice of the developer internally—capturing feedback from community and events, translating friction points into actionable priorities for Product/Engineering and content updates.
- Built and scaled high-impact technical content and education, including producing/hosting the Docker Build show and delivering practical demos and tutorials that simplified complex workflows into repeatable developer patterns.
- Led flagship developer events and programs (e.g., DockerCon) to drive product adoption and community trust; created frameworks for talks, demos, and technical storytelling at scale.
- Revitalized and scaled strategic community advocacy programs (Docker Captains and Community Leaders) to increase developer advocacy, feedback quality, and peer-to-peer enablement.

Senior Engineering Manager
5/2017 - 3/2020

- Managed all product delivery schedules, specification reviews, planning, budgeting, and resource allocation for: docs.docker.com, success.docker.com, support.docker.com
- Responsible for driving the architectural direction for Docker's Enterprise Customer Portal.
- Led the requirements gathering, planning, development, budgeting, and delivery of Docker's self-paced online learning platform.
- Lead the re-architecture of four separate monolithic properties into one unified micro services architecture using CI/CD, monitoring, AWS and Docker Swarm.
- Drove infrastructure-as-code adoption using Terraform to manage AWS resources across Docker's customer-facing properties, standardizing environment provisioning and reducing manual configuration drift.
- Managed 15 person team across 3 timezones including contractors and vendors.
- Technologies used Node.js, React, Redux, Express, Restify, MongoDb, Reddis, Elasticsearch, Docker, Kubernetes

2/2014-5/2017 | Supernaut
Director of Technology

- Responsible for overseeing and managing the technical direction of Supernaut's projects and deliverables.
- Manage a combined budget of 1.8M across Supernaut's projects including resource planning, hiring and road-mapping.
- Responsible for supporting business development efforts by scoping, drafting, estimating, and delivering proposals.
- Architected and shipped serverless backends on AWS Lambda and API Gateway for client engagements, designing event-driven workflows and REST/JSON APIs that reduced operational overhead and scaled with client demand.
- Architect and code full-stack solutions for our clients.
- Managed and mentored a team of Off-shore, Near-shore and On-shore project managers, software developers and testers.
- Primarily development in JavaScript, ReactJS, Redux, Angular, Node.js, MongoDb

8/2013-2/2014 | Click Security
Senior Front-End Engineer

- Senior member of the Applications development team designing and developing the UI for a network security application.
- The UI is developed using Angular.js.
- Designed and developed custom directives.
- Intimate knowledge of angular.js. Helped find and fix a memory leak in Angular.js framework.
- Develop middle tier modules in Node.js and MongoDB.